

Kernel Data Structures

Process Scheduling

Department of Computer Science
and Engineering



INDIAN INSTITUTE OF TECHNOLOGY
KHARAGPUR

Sandip Chakraborty
sandipc@cse.iitkgp.ac.in

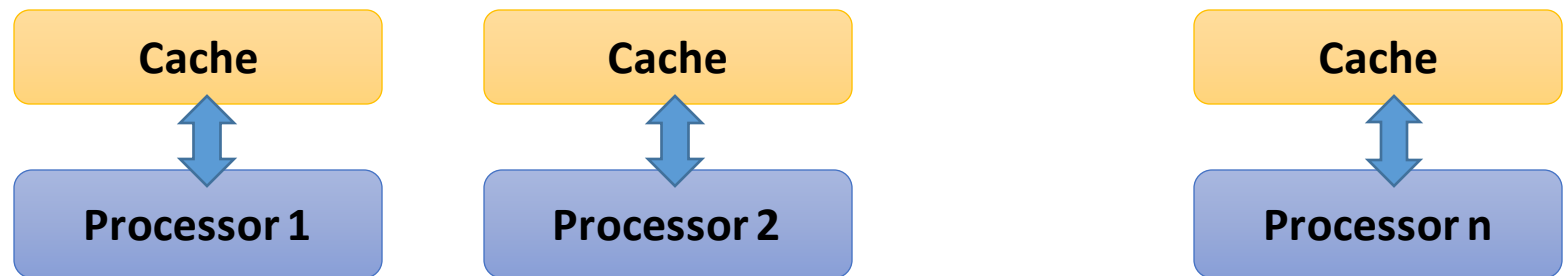


Linux Scheduling

- Linux is a multi-tasking Operating System
 - The system can execute multiple processes simultaneously

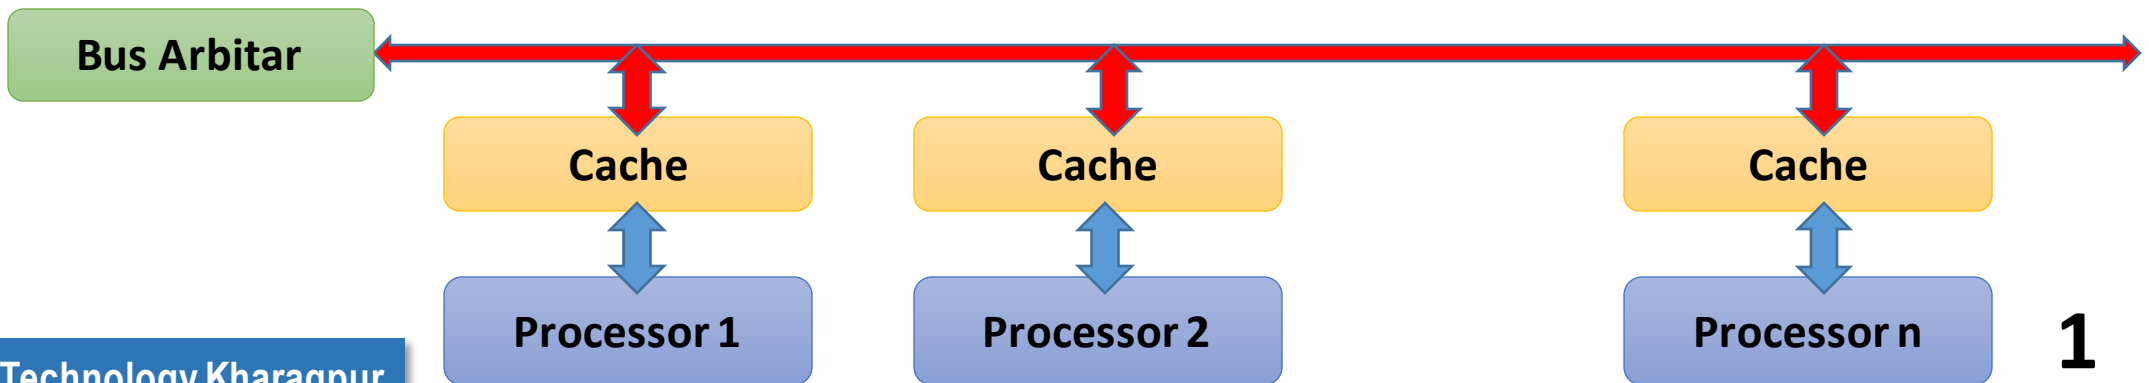
Linux Scheduling

- Linux is a multi-tasking Operating System
 - The system can execute multiple processes simultaneously
- **Symmetric Multiprocessing (SMP):**
 - Processing of programs by multiple processors that share a common operating system and memory
 - Linux kernel was designed to support SMP



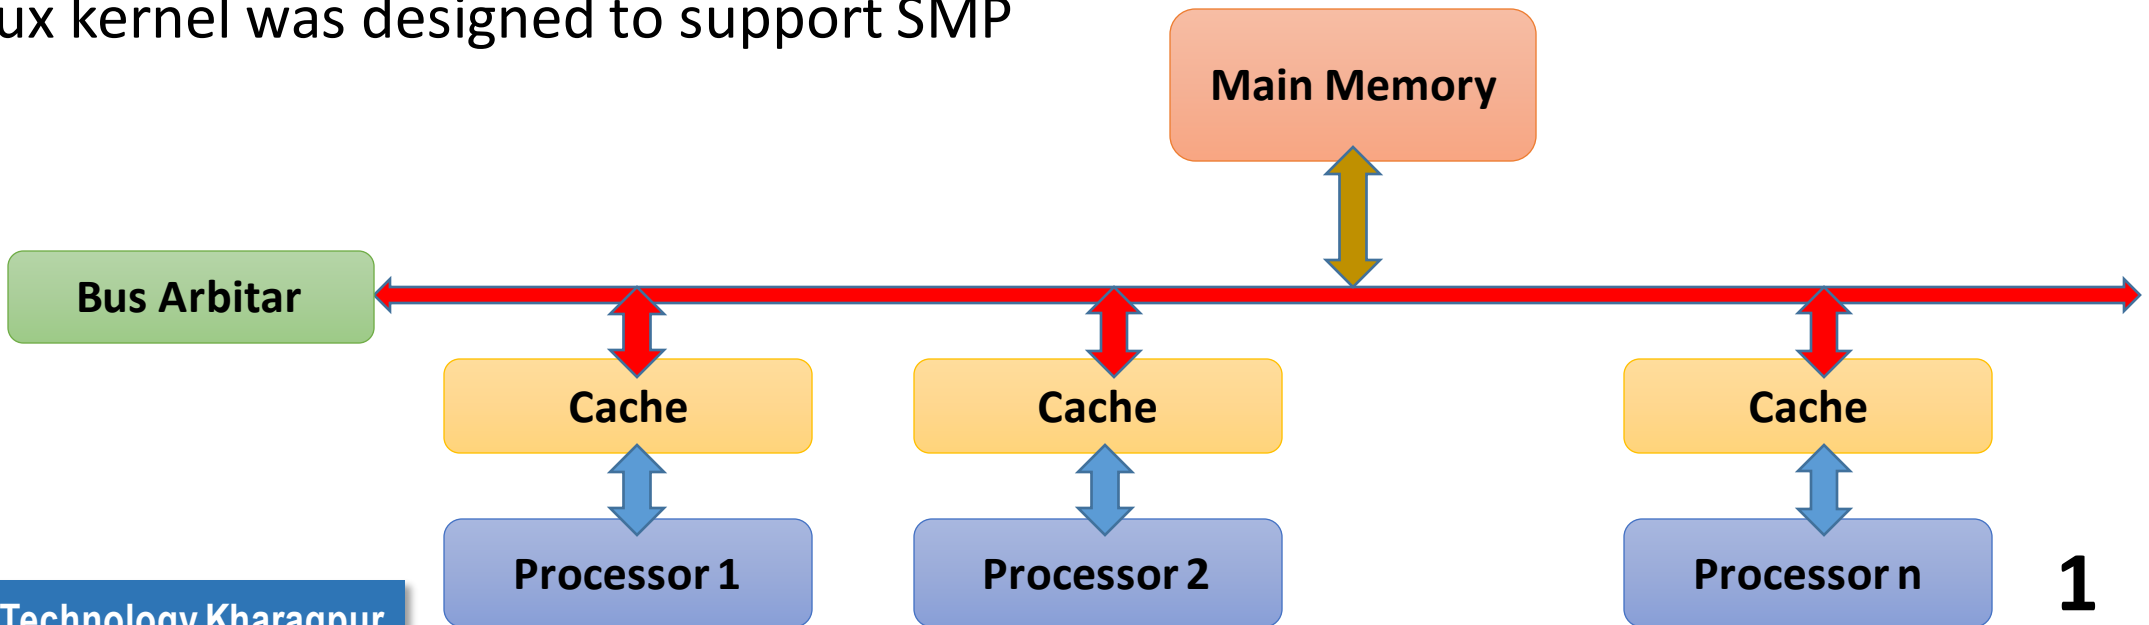
Linux Scheduling

- Linux is a multi-tasking Operating System
 - The system can execute multiple processes simultaneously
- **Symmetric Multiprocessing (SMP):**
 - Processing of programs by multiple processors that share a common operating system and memory
 - Linux kernel was designed to support SMP



Linux Scheduling


- Linux is a multi-tasking Operating System
 - The system can execute multiple processes simultaneously
- **Symmetric Multiprocessing (SMP):**
 - Processing of programs by multiple processors that share a common operating system and memory
 - Linux kernel was designed to support SMP




Linux Scheduling



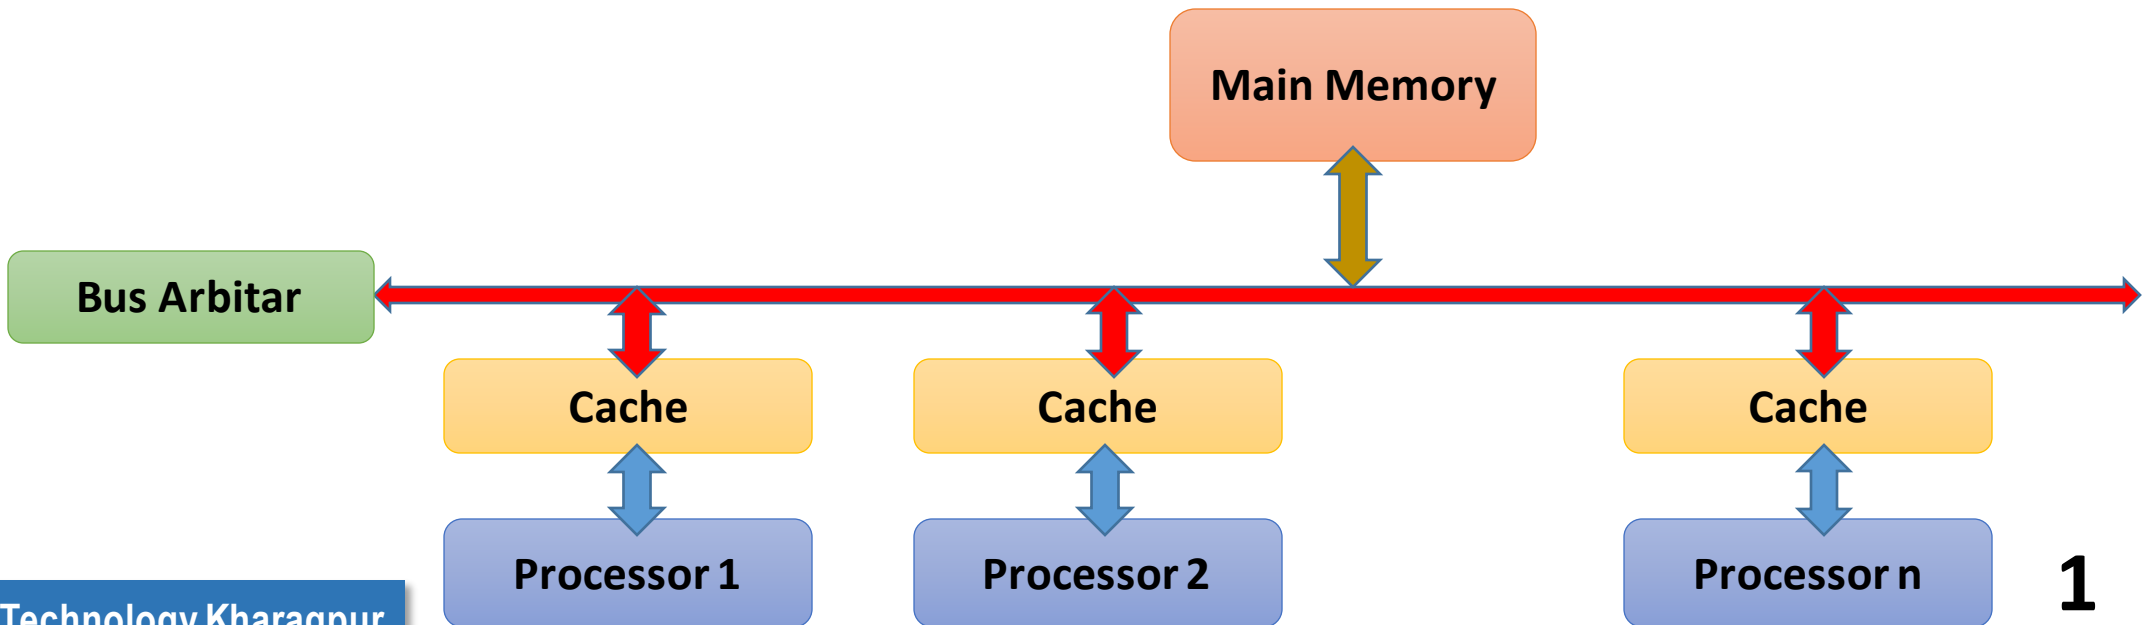
How do you assign
processes to processors?



Linux Scheduling



How do you assign processes to processors when the **different process has different priority** ?



Looking Back at the Process Descriptor

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```


Looking Back at the Process Descriptor

Task priority

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```

Looking Back at the Process Descriptor

Real-time task priority

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```



Looking Back at the Process Descriptor

Real-time task priority

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```

Looking Back at the Process Descriptor

Scheduling Class:

A generic structure to implement various scheduling algorithms

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```



sched_class

```
struct sched_class {  
    const struct sched_class *next;  
  
    void (*enqueue_task) (struct rq *rq, struct task_struct *p, int flags);  
    void (*dequeue_task) (struct rq *rq, struct task_struct *p, int flags);  
    void (*yield_task) (struct rq *rq);  
    bool (*yield_to_task) (struct rq *rq, struct task_struct *p,  
                           bool preempt);  
    void (*check_preempt_curr) (struct rq *rq, struct task_struct *p,  
                               int flags);  
  
    ...  
}
```

Looking Back at the Process Descriptor

Scheduling Entity:

Used for group
scheduling

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```



Looking Back at the Process Descriptor

Scheduling Entity (for RT tasks):

Used for group
scheduling for real-time
tasks

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```

Looking Back at the Process Descriptor

Scheduling Policy:

SCHED_NORMAL
SCHED_BATCH
SCHED_IDLE
SCHED_FIFO or
SCHED_RR

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```


Looking Back at the Process Descriptor

SCHED_NORMAL

Scheduling policy for regular tasks

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```

Looking Back at the Process Descriptor

SCHED_BATCH

Scheduling policy for batch jobs – does not preempt often as regular tasks

Ex: Bulk database updates

Allow tasks to run longer and make better use of caches

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```

Looking Back at the Process Descriptor

SCHED_IDLE

Scheduling policy for very low priority processes, like background tasks

Objective is not to disturb the regular tasks

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```

Looking Back at the Process Descriptor

SCHED_FIFO

SCHED_RR

Scheduling policy for real time processes

Handled by real-time schedulers

kernel/sched/rt.c

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```

Looking Back at the Process Descriptor

A bitmask indicating a task's affinity towards a CPU


Better utilization of the CPU cache

```
struct task_struct {  
    ...  
    int prio, static_prio, normal_prio;  
    unsigned int rt_priority;  
    const struct sched_class *sched_class;  
    struct sched_entity se;  
    struct sched_rt_entity rt;  
    ...  
    unsigned int policy;  
    cpumask_t cpus_allowed;  
    ...  
};
```

Looking Back at the Process Descriptor

A bitmask indicating a task's affinity towards a CPU

Better utilization of the CPU cache – **remember the SMP architecture**

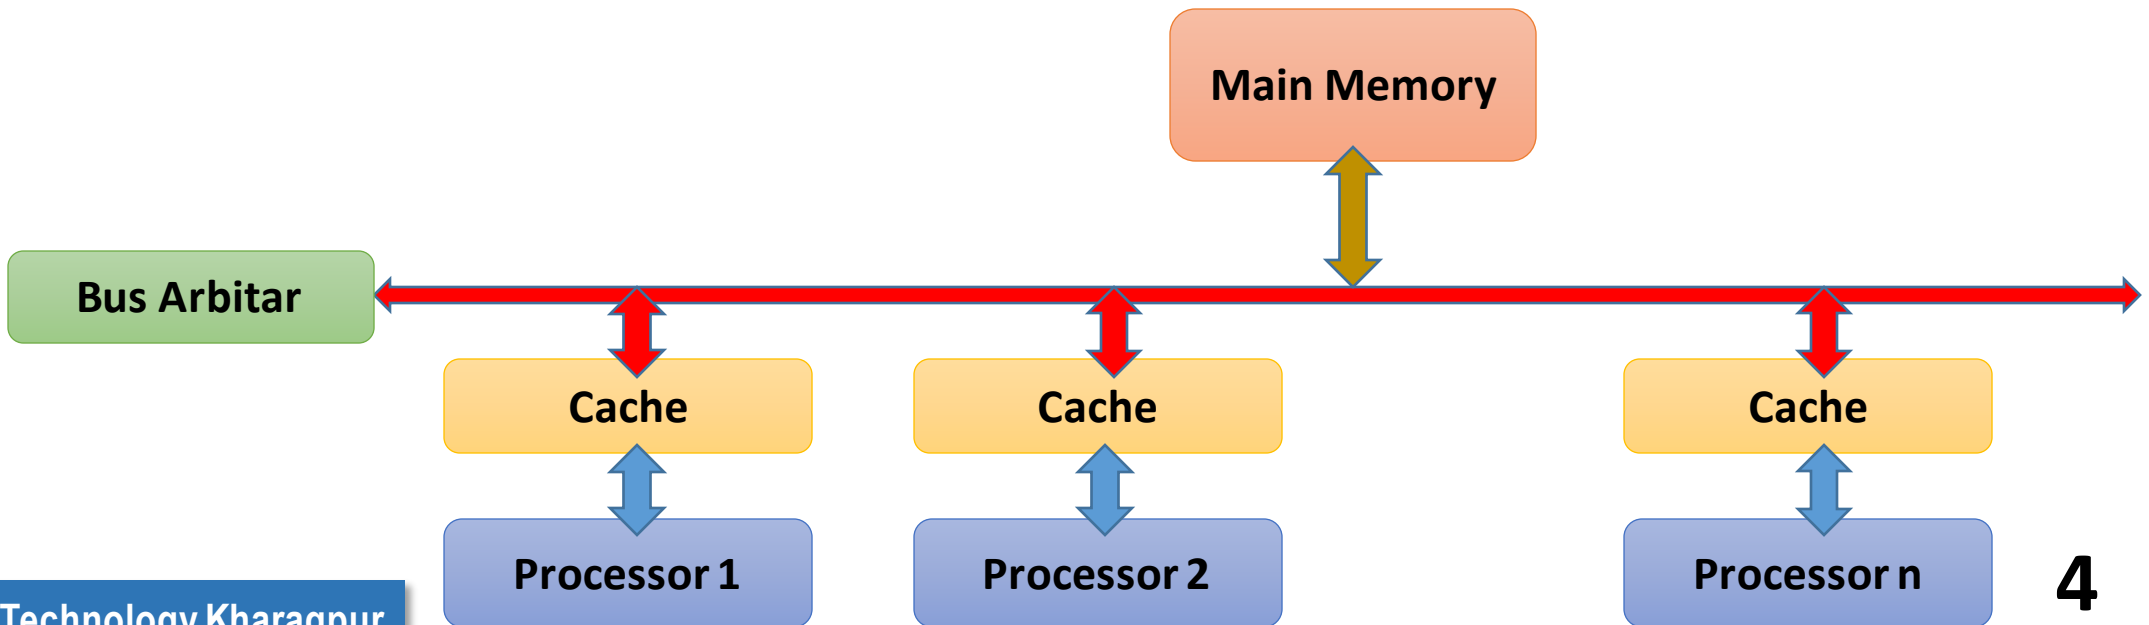


Looking Back at the Process Descriptor

A bitmask indicating a task's affinity towards a CPU

Better utilization of the CPU cache – **remember the SMP architecture**

Defines the **CPU Affinity** of a process !!



Process Priority

- Any UNIX-based system (hence, Linux) implements a priority-based scheduling
 - Assign some value to every task – indicates how important this task is compared to other tasks in the system

Process Priority

- Any UNIX-based system (hence, Linux) implements a priority-based scheduling
 - Assign some value to every task – indicates how important this task is compared to other tasks in the system
- Each task is assigned a nice value
 - An integer between –20 to 19 with default being 0

Process Priority

- Any UNIX-based system (hence, Linux) implements a priority-based scheduling
 - Assign some value to every task – indicates how important this task is compared to other tasks in the system
- Each task is assigned a nice value
 - An integer between –20 to 19 with default being 0
 - The higher the niceness – the lower the priority (**it is "nice" to other tasks**)
 - Use `<PS -Al>` (lowercase L) to check the niceness of the processes running in your system

Process Priority

- Use kernel system call `nice (int increment)` to change the niceness of a process

Process Priority

- Use kernel system call `nice (int increment)` to change the niceness of a process
- Change niceness from the user-space
 - `nice -n increment process_name`
 - `renice -n priority -p PID`


Real-time Tasks

- Execution should complete within a time boundary
- **Hard Real-time**
 - Strict time limits – the tasks must be completed within this time limit
 - By default, Linux does not support hard real-time processes
- **Soft Real-time**
 - The scheduler tries its best to maintain the time limit
 - However, the process can run a bit late depending on the current load

Process Priority (Kernel's Perspective)

- Kernel sees priorities in a different way than user
 - User may set the priority, but the kernel decides what to do with that priority
 - There are processes which are not under the control of any user-space program

Process Priority (Kernel's Perspective)

- Kernel sees priorities in a different way than user
 - User may set the priority, but the kernel decides what to do with that priority
 - There are processes which are not under the control of any user-space program
- Kernel uses a scale from 0 to 139 to represent the task priority values
 - 0 to 99 are reserved for real-time (soft) processes
 - 100 to 139 (mapped to nice -20 to +19) are for normal processes 

Process Priority (Kernel's Perspective)

- Kernel sees priorities in a different way than user
 - User may set the priority, but the kernel decides what to do with that priority
 - There are processes which are not under the control of any user-space program
- Kernel uses a scale from 0 to 139 to represent the task priority values
 - 0 to 99 are reserved for real-time (soft) processes
 - 100 to 139 (mapped to nice -20 to $+19$) are for normal processes
- Kernel sets the priority of a task depending on different factors
 - Real-time vs normal tasks
 - `static_prio` set by the user-space program (through nice)
 - The scheduling policy being used

The `schedule()` Function

- Main entry point to the kernel task scheduler
 - Replace the currently running task with a new task – **Context Switch**

The `schedule()` Function

- Main entry point to the kernel task scheduler
 - Replace the currently running task with a new task – **Context Switch**
- **But, what is a scheduler?**
 - A process?
 - A hardware module?
 - Something else?

The schedule() Function

- Main entry point to the kernel task scheduler
 - Replace the currently running task with a new task – **Context Switch**
- **But, what is a scheduler?**
 - A process?
 - A hardware module?
 - Something else?
- **Who does schedule the scheduler?**



Let's try to have a simplistic view ...

P1



Let's try to have a simplistic view ...

P1



Let's try to have a simplistic view ...

P1



`fork()`

P2




Let's try to have a simplistic view ...

P1



`fork()`

P2



Runqueue



Let's try to have a simplistic view ...

P1



fork()

P2



Runqueue



Let's try to have a simplistic view ...

P1





fork()

P2



schedule()

Runqueue



Let's try to have a simplistic view ...

P1




fork()

P2



schedule()

Runqueue



Runqueue gets updated, context switching is being done



Let's try to have a simplistic view ...


P1



P2



Runqueue




Let's try to have a simplistic view ...

`schedule()`

P1



P2



Runqueue



Let's try to have a simplistic view ...

`schedule()`

P1



P2



Runqueue



Let's try to have a simplistic view ...

P1



P2



fork()

Runqueue



Let's try to have a simplistic view ...

P1



P2



fork()

P3 (User)



Runqueue



Let's try to have a simplistic view ...

P1



P2




fork()

P3 (User)



Runqueue



Let's try to have a simplistic view ...

P1



P2





schedule()

P3 (User)



Runqueue



Let's try to have a simplistic view ...

P1



P2




schedule()


P3 (User)




Runqueue




Let's try to have a simplistic view ...




Runqueue



Let's try to have a simplistic view ...



Runqueue



Let's try to have a simplistic view ...

P1




P2



P3 (User)



Runqueue



Let's try to have a simplistic view ...

P1




P2



P3 (User)



Runqueue



How do you control the execution of a user process?



Let's try to have a simplistic view ...

P1




P2



P3 (User)




Runqueue



How do you control the execution of a user process?

Hardware to rescue !!!



Let's try to have a simplistic view ...

P1




P2




P3 (User)



Runqueue



Timer Interrupt



Let's try to have a simplistic view ...

P1




P2




P3 (User)



Runqueue



Timer Interrupt




Timer Interrupt
Handler
`/kernel/timer.c`

Let's try to have a simplistic view ...

P1




P2




P3 (User)



Runqueue



Timer Interrupt



Timer Interrupt
Handler

/kernel/timer.c


```
update_process_t  
imes()
```

Let's try to have a simplistic view ...

P1




P2




P3 (User)



Runqueue



Timer Interrupt



Timer Interrupt
Handler

/kernel/timer.c

update_process_t
imes()


scheduler_tick()

Let's try to have a simplistic view ...

P1



P2



P3 (User)



Runqueue




`scheduler_tick()`



Check the current
runqueue



Timer Interrupt



Timer Interrupt
Handler

`/kernel/timer.c`

`update_process_t
imes()`

`scheduler_tick()`

Let's try to have a simplistic view ...

P1



P2



P3 (User)



Runqueue




`scheduler_tick()`



Timer expires or high
priority process
waiting – Pre-empt



Timer Interrupt



Timer Interrupt
Handler

`/kernel/timer.c`

`update_process_t
imes()`

`scheduler_tick()`

Let's try to have a simplistic view ...

P1



P2



P3 (User)



Runqueue




`scheduler_tick()`



Timer expires or high
priority process
waiting – Pre-empt



Timer Interrupt



Timer Interrupt
Handler

`/kernel/timer.c`

`update_process_t
imes()`

`scheduler_tick()`

Let's try to have a simplistic view ...

P1



P2



P3 (User)



Runqueue




`scheduler_tick()`



Timer expires or high
priority process
waiting – Pre-empt



Timer Interrupt



Timer Interrupt
Handler

`/kernel/timer.c`

`update_process_t
imes()`


`scheduler_tick()`

Let's try to have a simplistic view ...

P1



P2



P3 (User)



Runqueue




`scheduler_tick()`



Timer expires or high
priority process
waiting – Pre-empt



Timer Interrupt



Timer Interrupt
Handler

`/kernel/timer.c`

`update_process_t
imes()`


`scheduler_tick()`

Let's try to have a simplistic view ...

P1



P2



P3 (User)



Runqueue



`scheduler_tick()`




Need for preemption



Call `schedule()`

Timer Interrupt



Timer Interrupt
Handler

`/kernel/timer.c`

`update_process_t
imes()`


`scheduler_tick()`

Let's try to have a simplistic view ...

P1



P2



P3 (User)



Runqueue



`scheduler_tick()`




Need for preemption



Call `schedule()`

Timer Interrupt



Timer Interrupt
Handler

`/kernel/timer.c`

`update_process_t
imes()`

`scheduler_tick()`

Let's try to have a simplistic view ...

P1



P2



P3 (User)



The timer is triggered through a hardware clock, so it is applied on all the processes running – the currently running process will be preempted to run the Timer Interrupt Handler

Runqueue



`scheduler_tick()`




Need for preemption



Call `schedule()`

Timer Interrupt



Timer Interrupt Handler

`/kernel/timer.c`

`update_process_t
imes()`

`scheduler_tick()`

Let's try to have a simplistic view ...

P1



P2



P3 (User)




A (kernel) process can directly call the `schedule()` function or it is called through the function `scheduler_tick()`

Runqueue



`scheduler_tick()`




Need for preemption



`Call schedule()`

Timer Interrupt



Timer Interrupt Handler

`/kernel/timer.c`


`update_process_times()`

`scheduler_tick()`

Invoking the Scheduler

- The main scheduler function `schedule()` is called from many different places in the kernel and for various occasions.
 - The invocations can be direct or lazy
 - Lazy invocation does not call the function by its name – gives the kernel a hint that the scheduler needs to be called soon

Invoking the Scheduler

- The main scheduler function `schedule()` is called from many different places in the kernel and for various occasions.
 - The invocations can be direct or lazy
 - Lazy invocation does not call the function by its name – gives the kernel a hint that the scheduler needs to be called soon
- **Periodic Scheduler**
 - `scheduler_tick()` is called on every timer interrupt – frequency is set during the Kernel compilation time (default 1000)
 - Tells the kernel whether and when the process needs to be scheduled next depending on the possibility of preemption
 - In case an interrupt recommends  that this process needs to be preempted, then the same is done during the execution of `scheduler_tick()`

Invoking the Scheduler

- **Currently running task enters the sleep state**
 - The task voluntarily gives up the CPU, waits for certain event to happen
 - The calling task adds itself to a *wait-queue* – sets itself as **TASK_INTERRUPTABLE** (can be interrupted when the event occurs) or **TASK_UNINTERRUPTABLE** (does not respond to an interrupt, periodically checks itself whether the event has occurred).
 - Call `schedu1e()` right before it goes to sleep

Invoking the Scheduler

- **Currently running task enters the sleep state**
 - The task voluntarily gives up the CPU, waits for certain event to happen
 - The calling task adds itself to a *wait-queue* – sets itself as **TASK_INTERRUPTABLE** (can be interrupted when the event occurs) or **TASK_UNINTERRUPTABLE** (does not respond to an interrupt, periodically checks itself whether the event has occurred).
 - Call `schedule()` right before it goes to sleep
- **Sleeping task wakes up**
 - `wake_up()` function is executed in the corresponding wait queue
 - The task is set to runnable and put back in the runqueue
 - If the task has higher priority than other tasks in the runqueue, **TIF_NEED_RESCHED** flag is set – **Lazy invocation** (kernel often checks for this flag)

