

CS61065: Theory And Applications of Blockchain

Consensus in Permissioned Settings

Department of Computer Science
and Engineering




INDIAN INSTITUTE OF TECHNOLOGY
KHARAGPUR

Sandip Chakraborty
sandipc@cse.iitkgp.ac.in

Shamik Sural
shamik@cse.iitkgp.ac.in

Permissioned Model

- A blockchain architecture where users are authenticated a priori
 - A Membership Service Provider (MSP) helps to obtain the chain membership
- Users know each other
 - However, users may not trust each other – Security and consensus are still required.
- Run blockchain among known and identified participants



Permissioned Model – Use Cases

- Particularly interesting for business applications – execute contracts among a closed set of participants
- Example: Provenance tracking of assets in a supply chain



Executing Contracts over a Closed Network

- **Smart Contracts:** “A self-executing contract in which the terms of the agreement between the buyer and the seller is directly written into the lines of code” - <http://www.scalablockchain.com/>
- **Agreement on a Smart Contract Execution:**
 - Store the contract on a blockchain
 - Once an event is triggered, execute the codes locally on each peer
 - Generate transactions as the output of the contract execution
 - The peers of the blockchain network validates the transaction, and the output is committed in the blockchain – may trigger the next event to execute the code further

Executing Contracts over a Closed Network

- **Smart Contracts:** “A self-executing contract in which the terms of the agreement between the buyer and the seller is directly written into the lines of code” - <http://www.scalablockchain.com/>
- **Agreement on a Smart Contract Execution:**
 - Store the contract on a blockchain
 - Once an event is triggered, execute the codes locally on each peer
 - Generate transactions as the output of the contract execution
 - The peers of the blockchain network validates the transaction and the output is committed in the blockchain – more code further

**Do we really need to execute
the code on each peer?**

**When does each peer execute the
code?**

Smart Contract Agreement as a State Machine Replication

- Execute contract at a subset of nodes, and ensure that the same state is propagated to all the nodes
 - Majority of the peers should agree on the state
 - **Validation:** Generate a "*proof*" that a peer has agreed on the "*state of execution*"

Smart Contract Agreement as a State Machine Replication

- Execute contract at a subset of nodes, and ensure that the same state is propagated to all the nodes
 - Majority of the peers should agree on the state
 - **Validation:** Generate a "*proof*" that a peer has agreed on the "*state of execution*"




How will we generate the proof?

Smart Contract Agreement as a State Machine Replication




- Execute contract at a subset of nodes, and ensure that the same state is propagated to all the nodes
 - Majority of the peers should agree on the state
 - **Validation:** Generate a "*proof*" that a peer has agreed on the "*state of execution*"
- **State Machine Replication:**
 - Represent the smart contract as a state machine – **Remember, any deterministically executable code can be represented as a state machine**

```
S1:  
while (moreGoods == 1)  
    DeliverGoods();  
S2:  
if (allOrderComplete == 0) goto S1;  
else {  
    S3:  
    printf("Goods transfer complete");  
}
```




State Machine Replication

Replicate the state machine
on multiple independent
servers





State Machine Replication




T1

T1:
I have delivered 100kg
potatoes





State Machine Replication

T1




T2:
I have delivered 50kg
tomatoes

T2





State Machine Replication

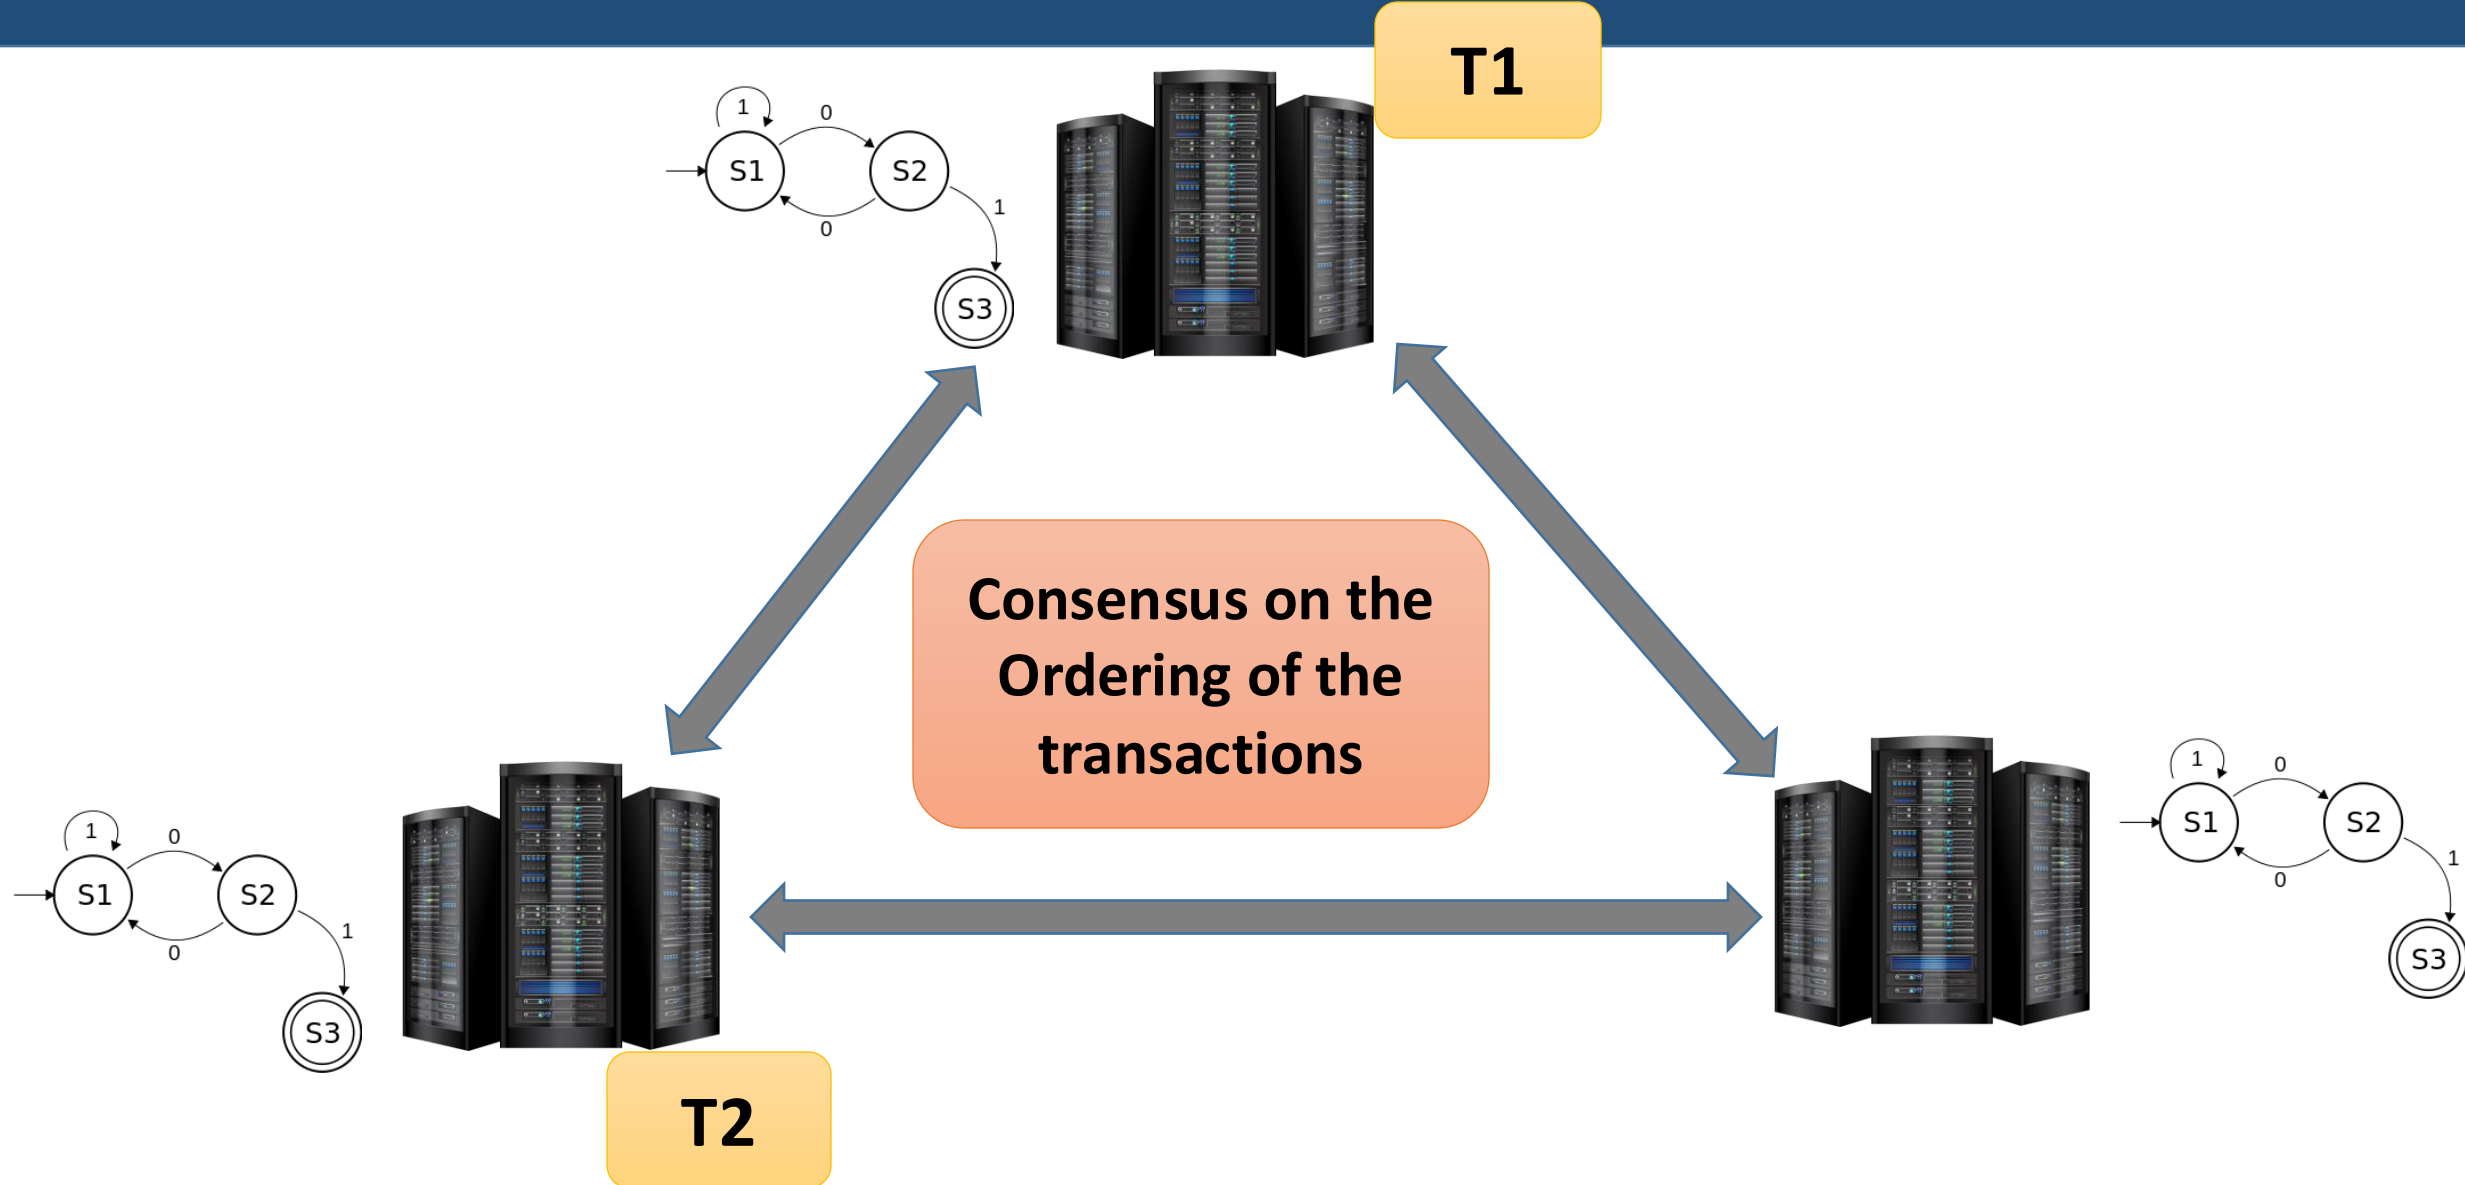
T1



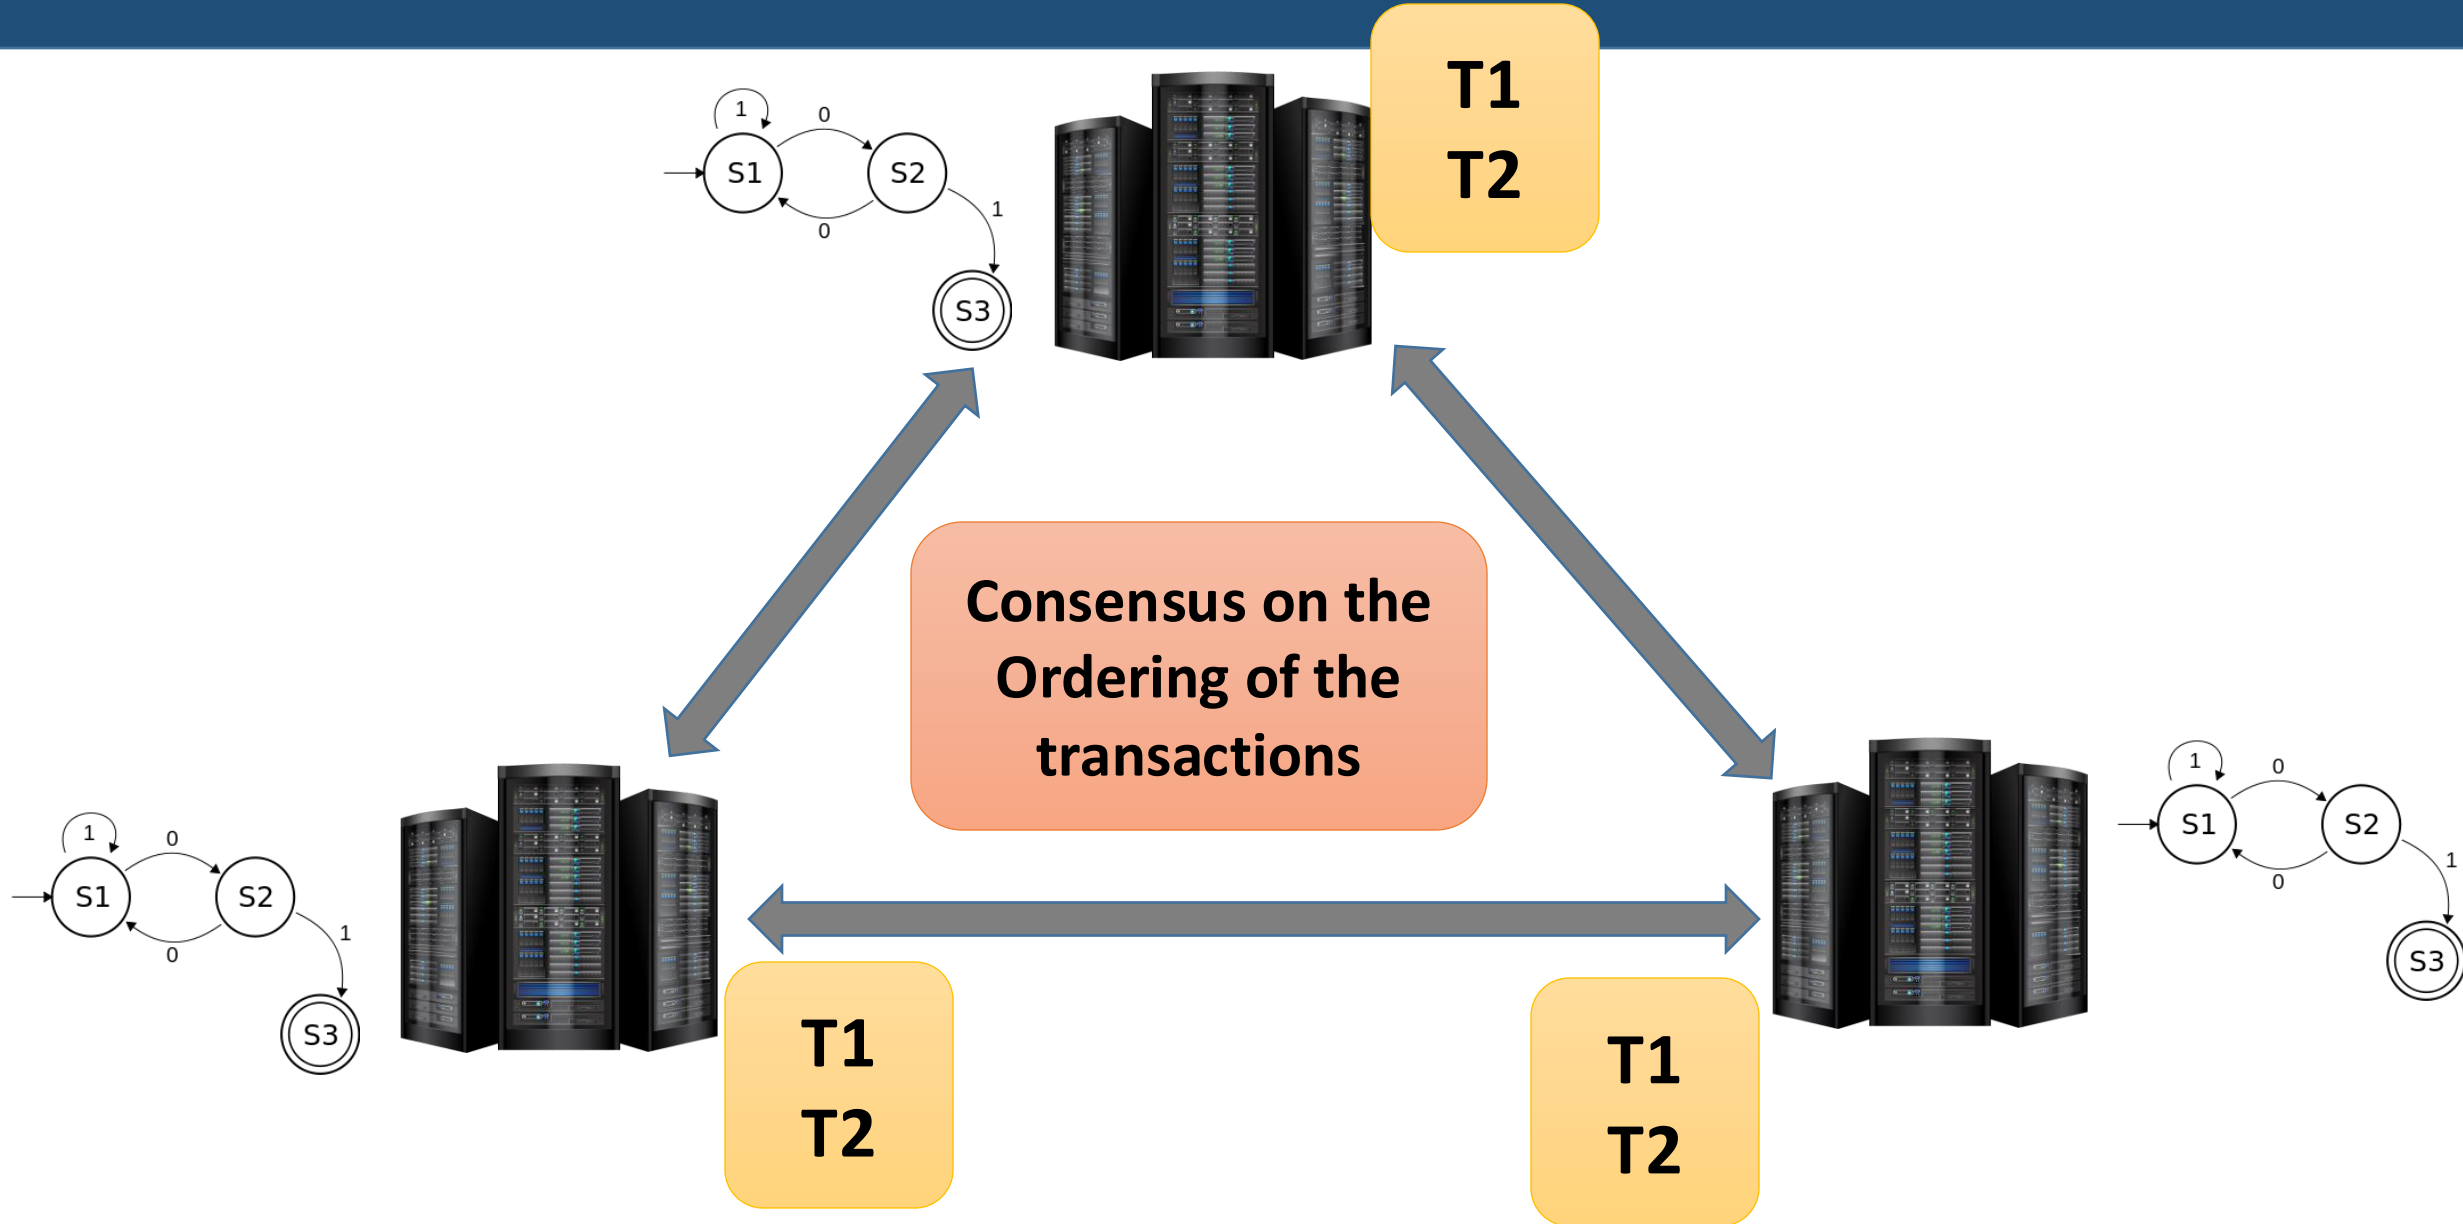
T2




State Machine Replication



State Machine Replication




State Machine Replication




T1
T2

**Independently
execute the
transactions**



T1
T2




T1
T2


State Machine Replication



Independently
execute the
transactions




T1
T2



T1
T2


State Machine Replication

More orders?
Yes




T1
T2

Independently
execute the
transactions




T1
T2



T1
T2


State Machine Replication

Execution of T1
completes



T1
T2


Independently
execute the
transactions



T1
T2




T1
T2




State Machine Replication

Start executing T2




T1
T2

Independently
execute the
transactions




T1
T2




T1
T2

State Machine Replication

Complete
execution




T1
T2



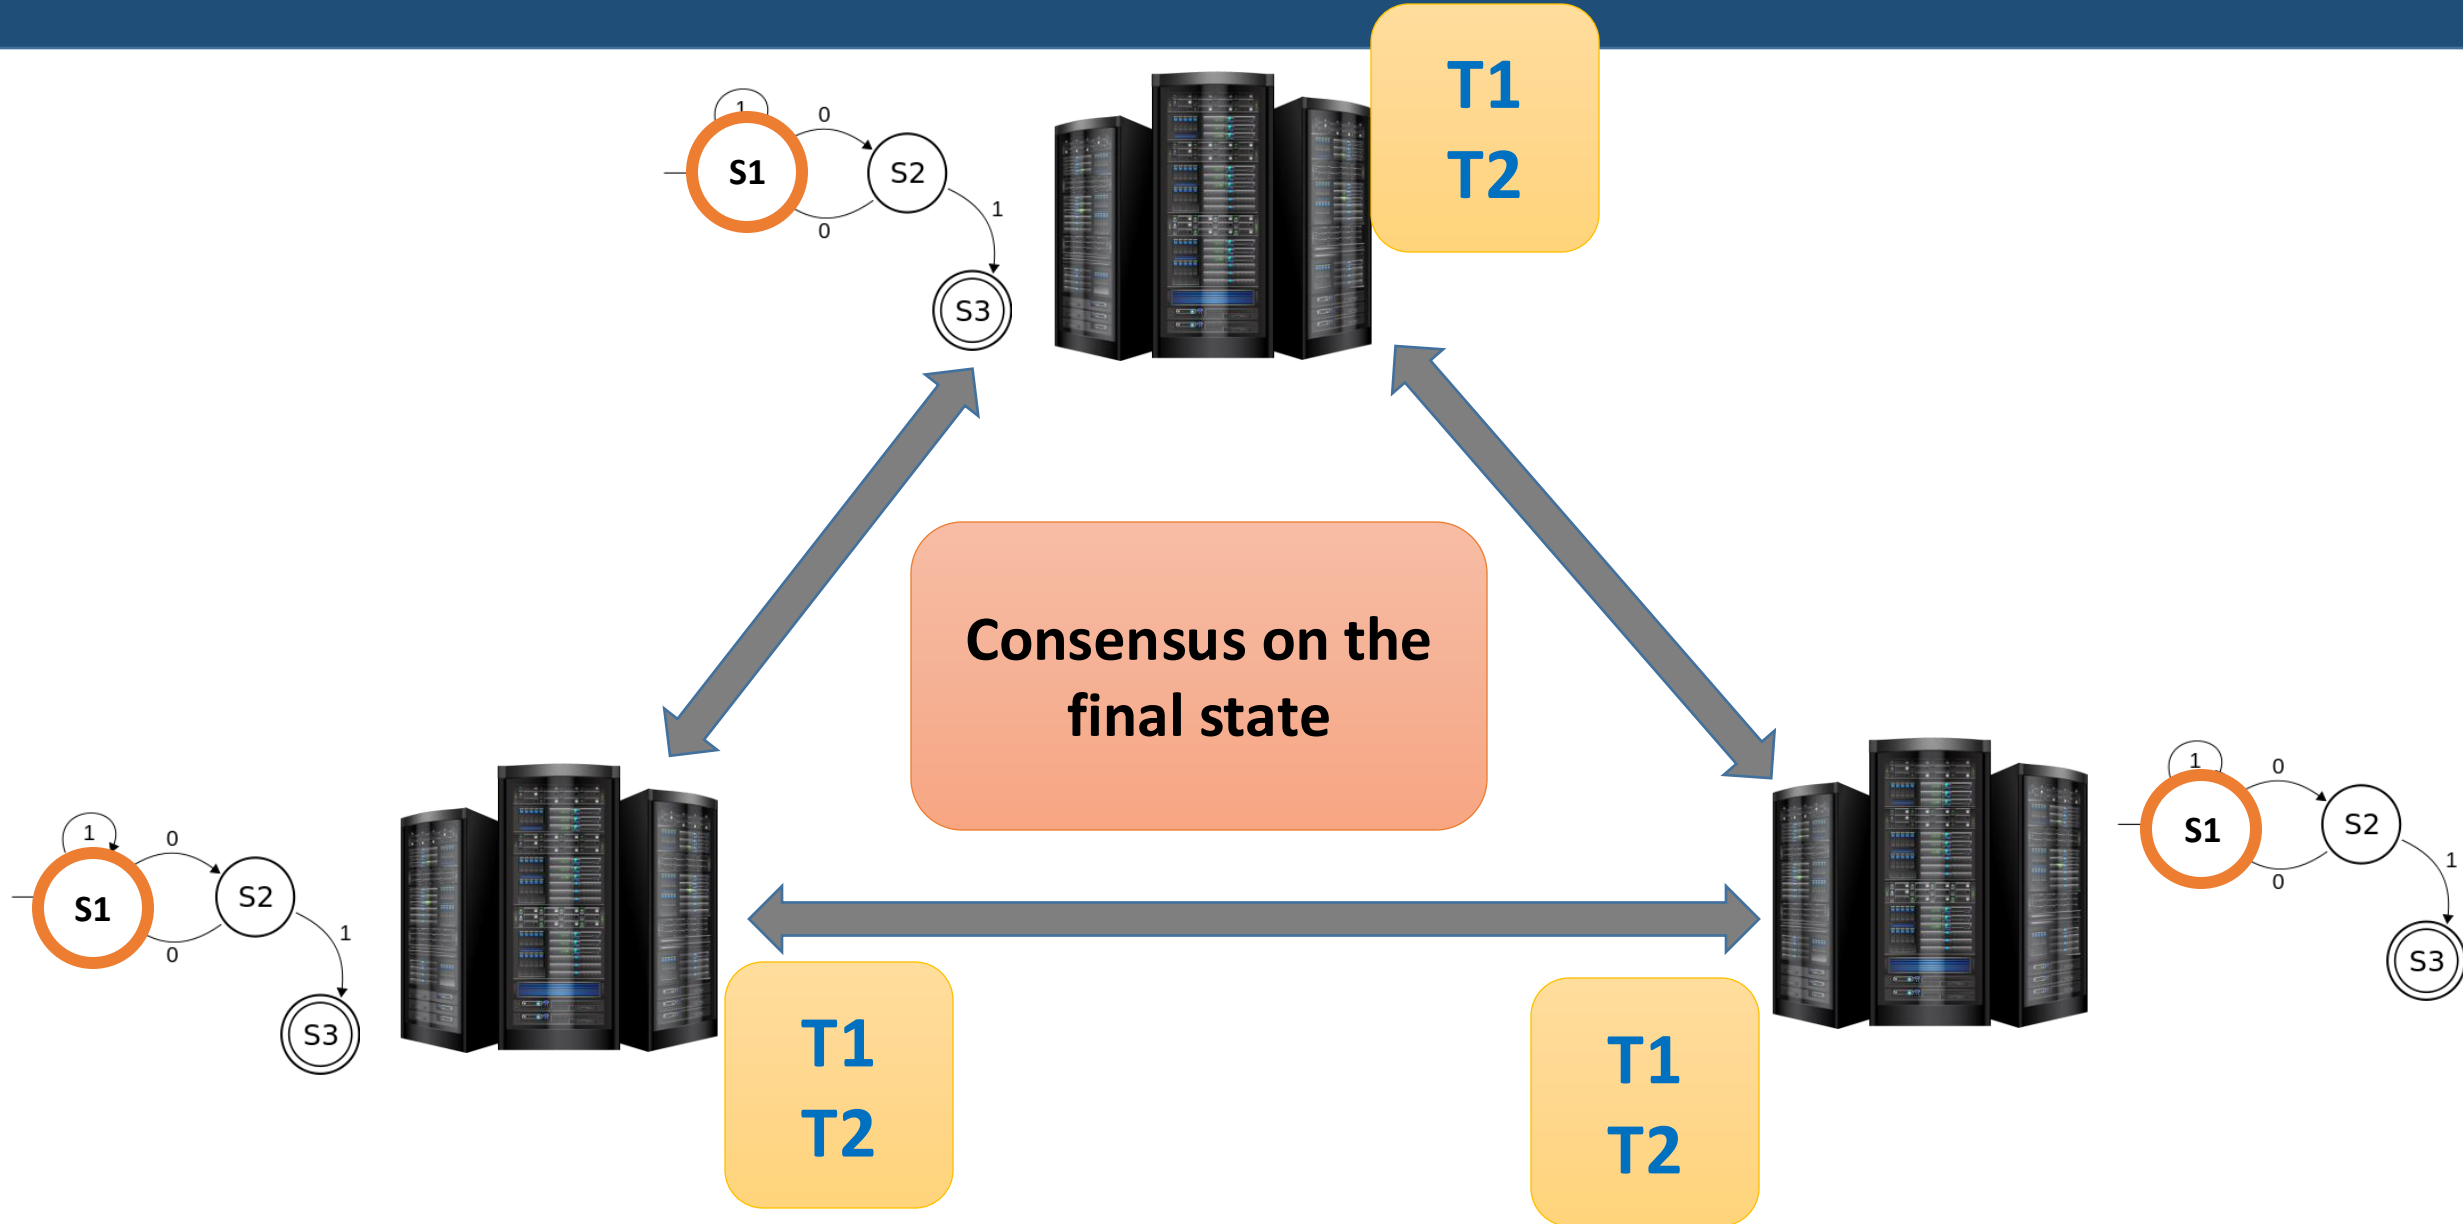
T1
T2




T1
T2



State Machine Replication




State Machine Replication



T1
T2




Inform



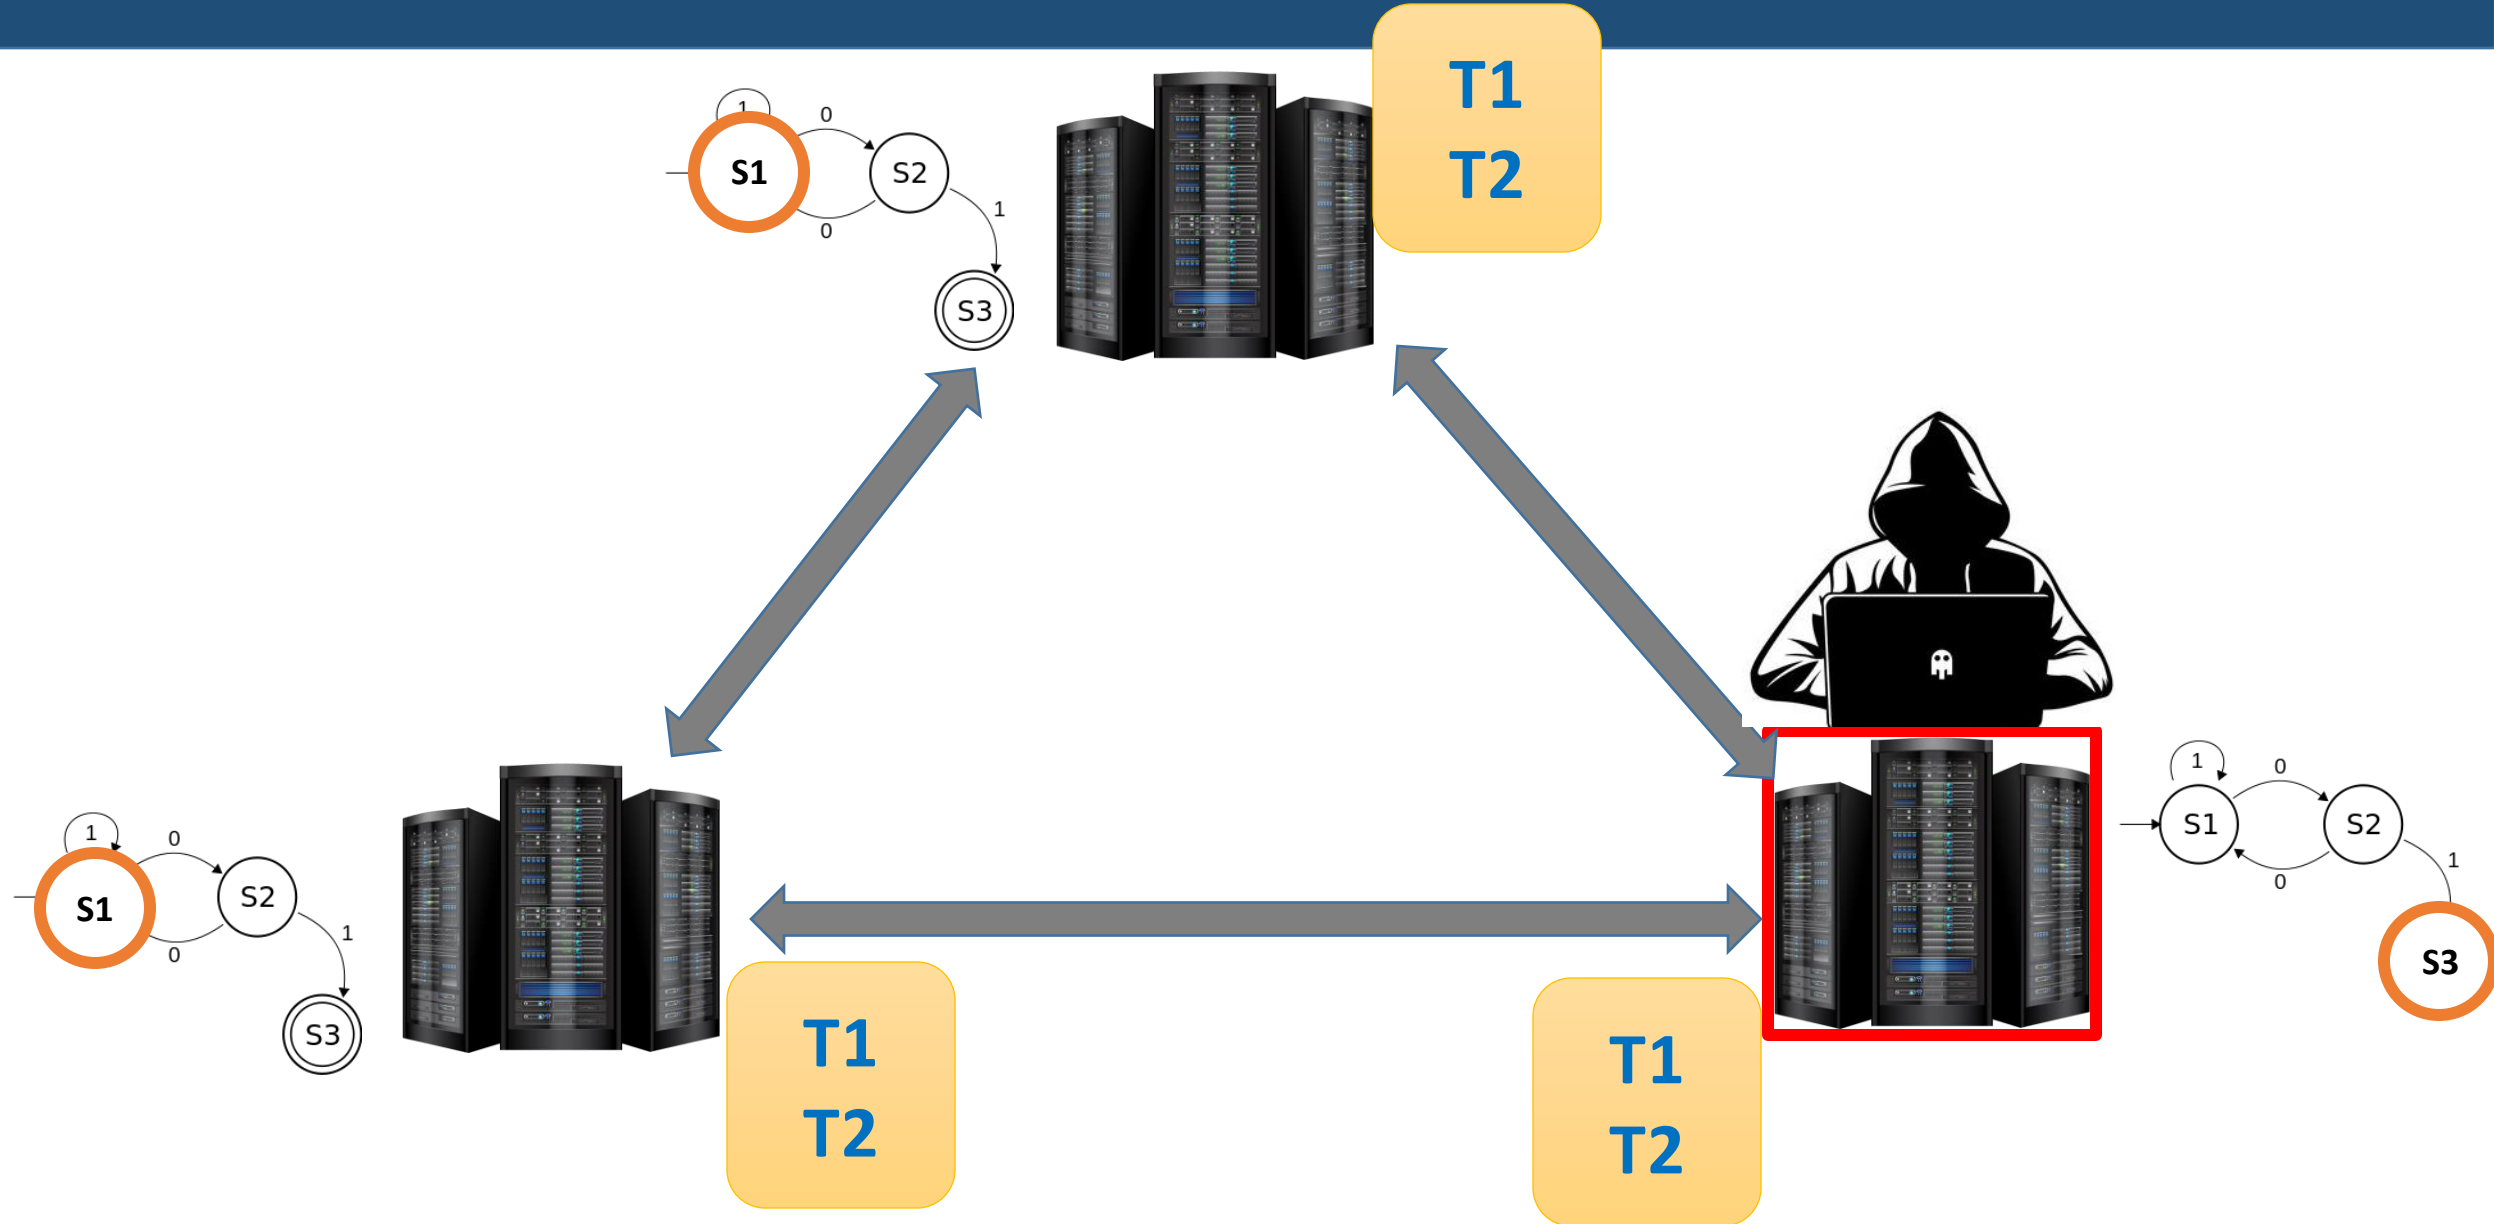
T1
T2



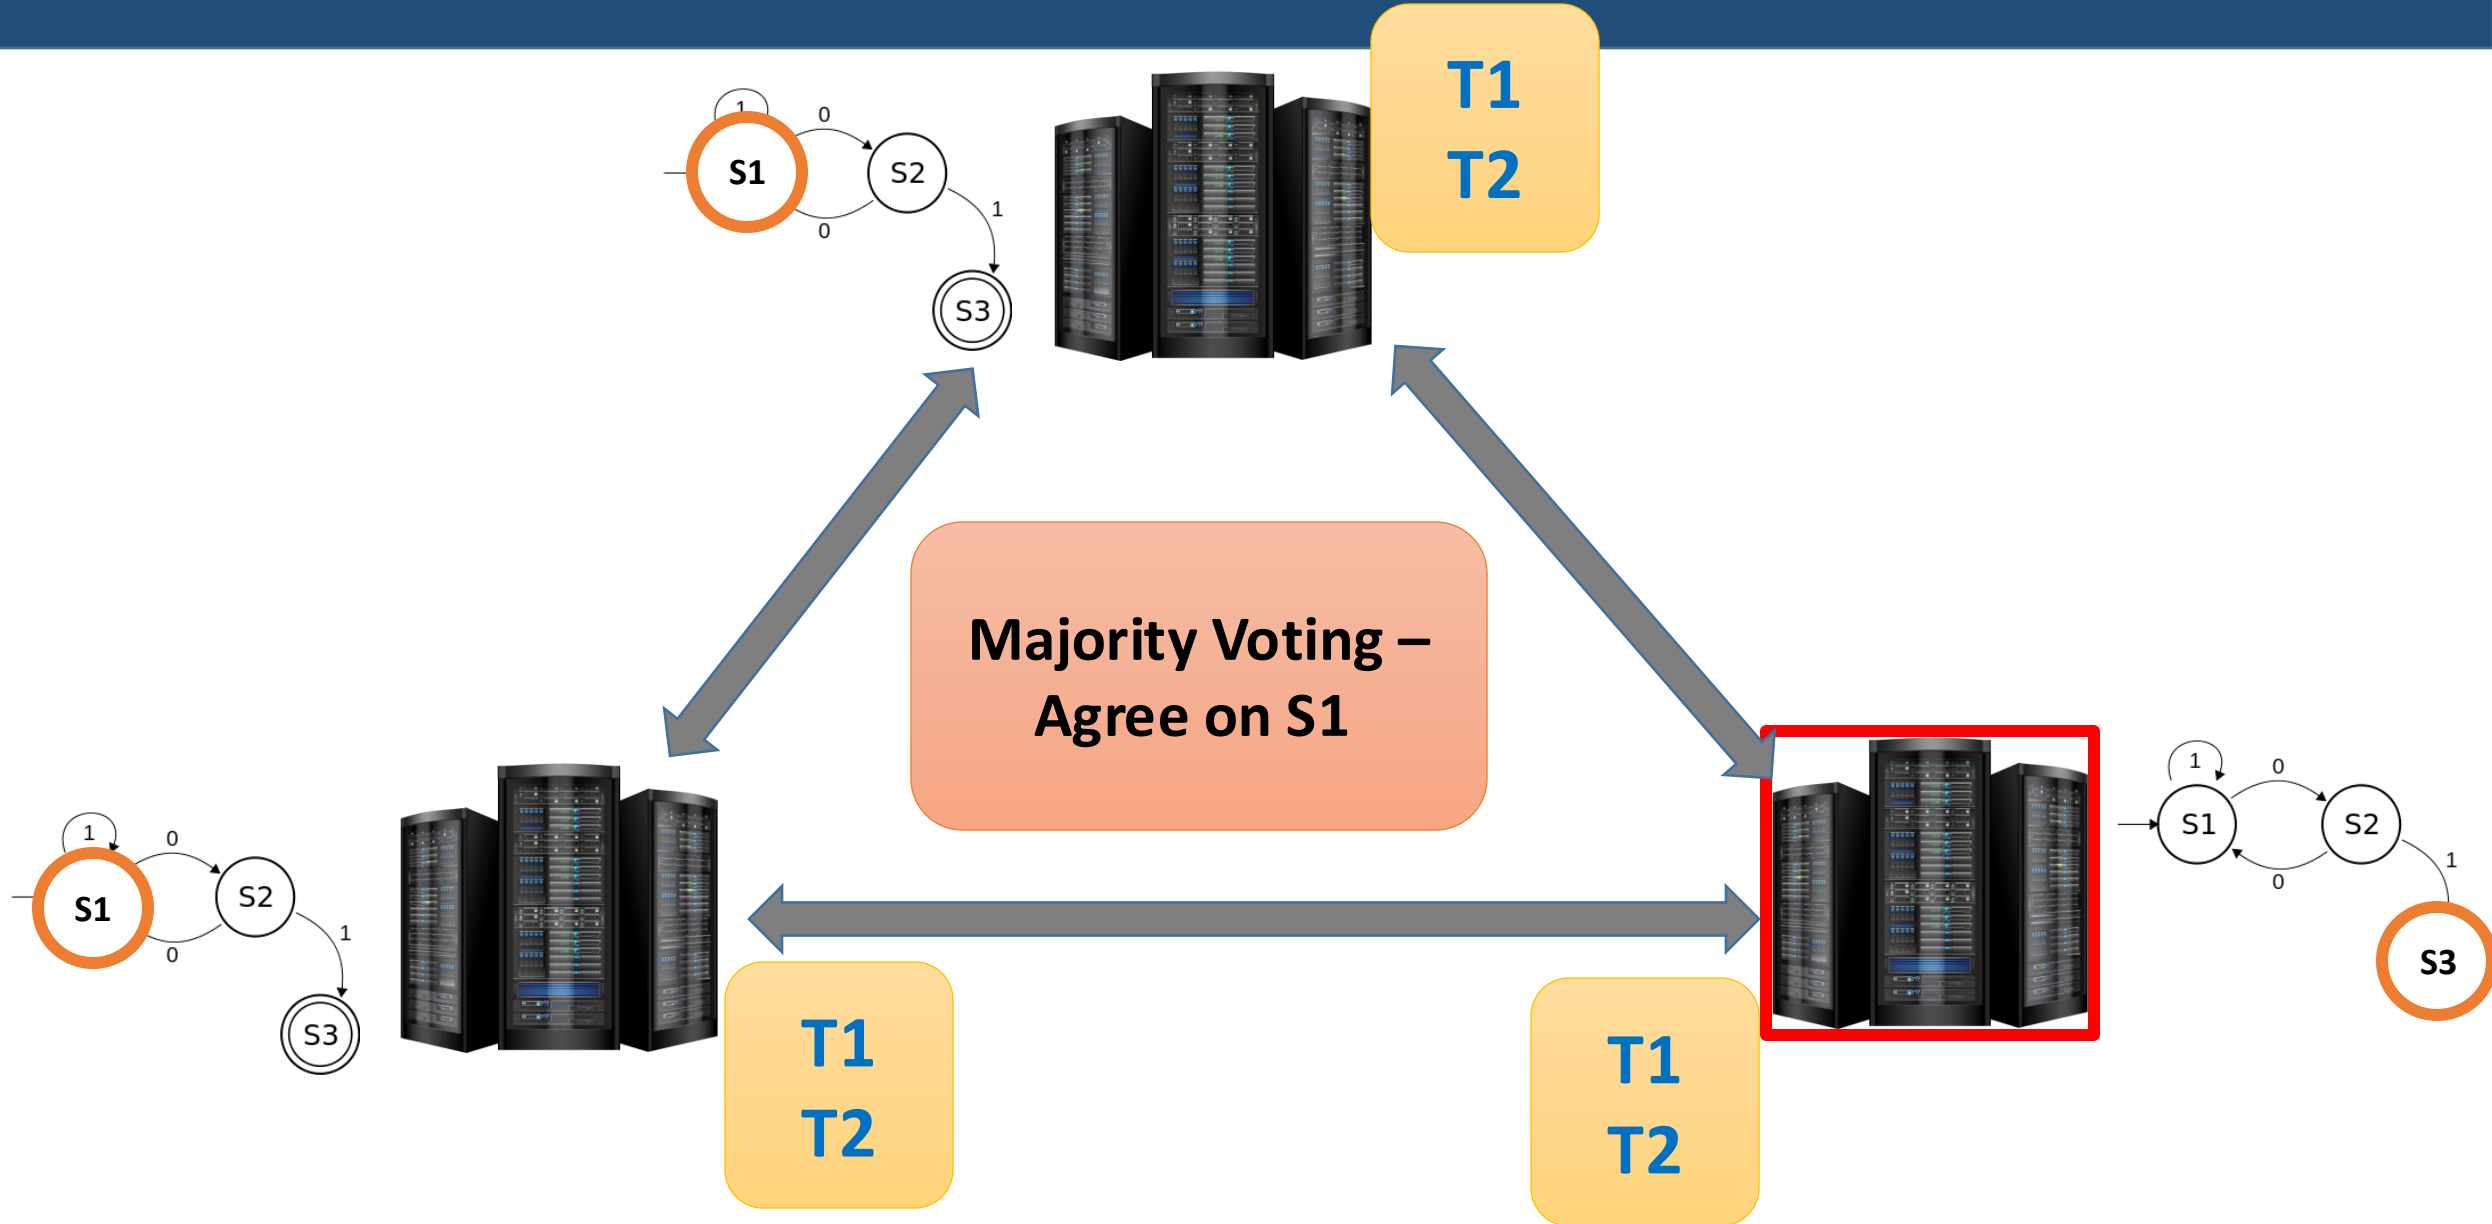
T1
T2




State Machine Replication – Why do we need Consensus?



State Machine Replication – Why do we need Consensus?




State Machine Replication – Why do we need Consensus?



T1
T2




What will happen
to this client?



T1
T2



T1
T2



Permissioned Blockchain and State Machine Replication

- There is a natural reason to use state machine replication-based consensus over permissioned blockchains
 - The network is closed, the nodes know each other, so state replication is possible among the known nodes
 - Avoid the overhead of mining - do not need to spend anything (like power, time, bitcoin) other than message passing
 - However, consensus is still required - machines can be faulty or behave maliciously

Permissioned Blockchain and State Machine Replication

- There is a natural way to implement a replication-based consensus over permissioned nodes
 - The network is known in advance, so state replication is possible among the nodes
 - Avoid the need for a proof of work (like bitcoin) or other resource-intensive mechanisms
 - Avoid the need to spend anything (like power, time, money)
 - However, nodes can be faulty or behave maliciously

But, we need a bit redesign !

Permissioned Blockchain and State Machine Replication

- There is a natural way to implement a replication-based consensus over permissioned nodes
 - The network is trusted, so state replication is possible among the nodes
 - Avoid the need for a trusted third party (like power, time, etc.)
 - However, nodes can be faulty or behave maliciously

But, we need a bit redesign !

Crypto is the saver

**Crypto + Distributed Consensus =
Consensus for Permissioned
Blockchain**

Permissioned Blockchain and State Machine Replication

- There is a natural reason to use state machine replication-based consensus over permissioned blockchains
 - The network is closed, the nodes know each other, so state replication is possible among the known nodes
 - Avoid the overhead of mining - do not need to spend anything (like power, time, bitcoin) other than message passing
 - However, consensus is still required - machines can be faulty or behave maliciously
- Classical Distributed Consensus Algorithms (Paxos, RAFT, Byzantine Agreement) are based on State Machine Replication
 - Let us (re)visit those algorithms

Faults in a Distributed System

- **Crash Faults:** The node stops operating – hardware or software faults
 - In an asynchronous system: You do not know whether messages have been delayed or the node is not responding
 - Rely on majority voting – progress as and when you have received the confirmation from the majority
 - Propagation of the consensus information – nodes on a slow network will receive it eventually

Faults in a Distributed System

- **Crash Faults:** The node stops operating – hardware or software faults
 - In an asynchronous system: You do not know whether messages have been delayed or the node is not responding
 - Rely on majority voting – progress as and when you have received the confirmation from the majority
 - Propagation of the consensus information – nodes on a slow network will receive it eventually
- **Byzantine Faults:** Nodes misbehave – send different information to different peers (partition the network)
 - More difficult to handle
 - More suitable for blockchains

Asynchronous Consensus with Crash Faults

- Remember the **FLP Impossibility**
 - Give priority to safety over liveness
- Guarantees the followings --
 - **Validity**: If all correct process proposes the same value v , then any correct process decides v
 - **Agreement**: No two correct processes decide differently
 - **Termination**: Every correct process eventually decides

Asynchronous Consensus with Crash Faults

- Remember the **FLP Impossibility**
 - Give priority to safety over liveness
- Guarantees the followings --
 - **Validity**: If all correct process proposes the same value v , then any correct process decides v (**Unlikely to happen in PoW**)
 - **Agreement**: No two correct processes decide differently (**Safety – Not in PoW**)
 - **Termination**: Every correct process eventually decides (**Liveness – Priority in PoW**)

Asynchronous Consensus with Crash Faults

- Remember the **FLP Impossibility**
 - Give priority to safety over liveness
- Guarantees the followings --
 - **Validity**: If all correct process proposes the same value v , then any correct process decides v
 - **Agreement**: No two correct processes decide differently
 - **Termination**: Every correct process eventually decides
- CFT Consensus
 - Paxos (Proposed by Lamport, the most fundamental CFT) -- used in DynamoDB
 - RAFT (Much simpler than Paxos) -- Used in **Fabric Transaction Ordering**

thank you!